

Ten questions to ask any AI vendor before an agent touches your ERP.

Hand this page to the vendor and ask which of these they implement, and how you can inspect it. "Yes" without evidence is a "no".

- | | |
|--|---|
| ☐ Read before write
Does the agent read the current SAP state before it proposes anything, and is the proposal diffed against that state? | ☐ Audit trail per case
Can every case be rebuilt from its sources, calculation snapshot, hypotheses, decisions and timestamps? |
| ☐ Confirm to act
Does a person review and confirm every create and update, with no path that skips the reviewer? | ☐ Explicit gaps
Does missing data stay visibly missing instead of becoming a zero or an empty success? |
| ☐ Verify after
Is SAP read back after every action and compared with the intent, with the result stored? | ☐ Untrusted text
Are documents and source comments treated as untrusted input, with scope and tool limits enforced outside the prompt? |
| ☐ Deterministic numbers
Do all financial figures come from versioned code rather than from the language model? | ☐ Credentials server-side
Do model credentials stay out of the browser, prompts, errors and audit records? |
| ☐ Bounded tools
Is the tool catalogue read-only by default, named and bounded, with no generic query or raw table access? | ☐ Evidence pack
Can InfoSec and internal audit receive tool policies, sample audit trails and read-back results before go-live? |

How SVLS LABS answers

All ten are implemented in our control model and demonstrated in Value Lens (private beta on synthetic SAP-like data; production SAP connector in development). We show the artefacts and the review log in a discovery call.

Book a discovery call: svslabs.com/contact/ · service@svslabs.com · +91 8500811119