Ten questions to ask any AI vendor before an agent touches your ERP.

Hand this page to the vendor and ask which of these they implement, and how you can inspect it. "Yes" without evidence is a "no".

- [ ] **Read before write**
  Does the agent read the current ERP state before it proposes anything, and is the proposal diffed against that state?
- [ ] **Audit trail per case**
  Can every case be rebuilt from its sources, calculation snapshot, hypotheses, decisions and timestamps?
- [ ] **Confirm to act**
  Does a person review and confirm every create and update, with no path that skips the reviewer?
- [ ] **Explicit gaps**
  Does missing data stay visibly missing instead of becoming a zero or an empty success?
- [ ] **Verify after**
  Is the ERP read back after every action and compared with the intent, with the result stored?
- [ ] **Untrusted text**
  Are documents and source comments treated as untrusted input, with scope and tool limits enforced outside the prompt?
- [ ] **Deterministic numbers**
  Do all financial figures come from versioned code rather than from the language model?
- [ ] **Credentials server-side**
  Do model credentials stay out of the browser, prompts, errors and audit records?
- [ ] **Bounded tools**
  Is the tool catalogue read-only by default, named and bounded, with no generic query or raw table access?
- [ ] **Evidence pack**
  Can InfoSec and internal audit receive tool policies, sample audit trails and read-back results before go-live?

How SVLS LABS answers

All ten are implemented in our control model and demonstrated in Value Lens (private beta on synthetic ERP-like data; production ERP connector in development). We show the artefacts and the review log in a discovery call.

Book a discovery call: svlslabs.com/contact/ · service@svlslabs.com · +91 8500811119